

Introduction to Quantum Machine Learning

Disclaimer: Work in progress. Portions of these written materials are incomplete.

Quantum Computing Basics

Lightning Recap

Origins of Quantum Computing

- Idea: making a computer out of quantum-mechanical building blocks in order to better mimic nature's quantum mechanical properties
- Richard Feynman (1982):
“Nature isn't classical, dammit, and if you want to make a simulation of nature, you'd better make it quantum mechanical.”

Quantum Mechanics Basics

- Quantum theory: a more general form of probability theory
 - Complex amplitudes allow for wave-like interference in the space of possibilities

$$p(x) \in \mathbb{R}^+$$

$$\psi(x) \in \mathbb{C}$$

Qubits as building blocks

Classical bit

$$p(0) + p(1) = 1$$

Quantum bit
(qubit)

$$\psi_0 |0\rangle + \psi_1 |1\rangle$$

$$\begin{pmatrix} \psi_0 \\ \psi_1 \end{pmatrix} \in \mathbb{C}^2$$

$$|\psi_0|^2 + |\psi_1|^2 = 1$$

Qubits as building blocks

$$\text{Single qubit} \cong \mathbb{C}^2$$

$$N \text{ qubits} \cong \mathbb{C}^{2^N} \cong \bigotimes_{j=1}^N \mathbb{C}^2$$

$$\psi_0 |0\rangle + \psi_1 |1\rangle$$

Superposition of 2^1
values (one bit)

$$\sum_{j \in \{0,1\}^n} \psi_j |j\rangle$$

Superposition of 2^N
values (bitstrings) $\mathcal{H} \cong \mathbb{C}^{2^N}$

Understanding quantum circuits

Quantum circuit: sequence of evolutions of a quantum memory state

gates

measurement

Registers
(qubits)

Outcome
data

Quantum Artificial Intelligence

Quantum AI: main areas of current interest

Quantum-enhanced optimization:

using quantum dynamics for an
optimization advantage

Quantum deep learning:

learning quantum representations
for quantum data

Quantum Deep Learning?

Classical vs. Quantum Deep Learning

- Parameters Φ ▪ Input x
- Feedforward operation f
- Loss function $L(f(x, \Phi), y)$
- Goal: find $\operatorname{argmin}_{\Phi} L(f(x), \Phi), y)$

- Parameters Φ ▪ Input $|\xi_0\rangle$
- Feedforward operation $\hat{U}(\Phi)$
- Loss operator \hat{L}
- Goal: find $\operatorname{argmin}_{\Phi} \langle \xi_{\Phi} | \hat{L} | \xi_{\Phi} \rangle$

Quantum-Classical Optimization of Quantum Neural Nets

QNN training loop:

1. Inference (sampling) done on QPU
2. Loss expectation estimate relayed to CPU
3. CPU suggests new parameters
4. Repeat 1-3 until convergence

Theory: Quantum Neural Nets & Gradients

1. QNN's general structure
2. Pauli operator expansion*
3. Quantum exponential compilation*
4. Quantum expectation estimation
5. Finite-difference gradients
6. Parameter-shift gradients*
7. Stochastic Gradient estimation*

*see [TFQ whitepaper](#) for detailed background

Theory: QNN's & Gradients

- Multi-layer QNN's: compilation hierarchy

$$\hat{U}(\boldsymbol{\theta}) = \prod_{\ell=1}^L \hat{V}^{\ell} \hat{U}^{\ell}(\boldsymbol{\theta}^{\ell}) \quad \hat{U}^{\ell}(\boldsymbol{\theta}^{\ell}) \equiv \bigotimes_{j=1}^{M_{\ell}} \hat{U}_j^{\ell}(\theta_j^{\ell})$$

$$\hat{U}_j^{\ell}(\theta_j^{\ell}) = e^{-i\theta_j^{\ell}\hat{g}_j^{\ell}}, \quad \hat{g}_j^{\ell} = \sum_{k=1}^{K_{j\ell}} \beta_k^{j\ell} \hat{P}_k,$$

$$\hat{U}_j^{\ell}(\theta_j^{\ell}) = \prod_k e^{-i\theta_j^{\ell}\beta_k^{j\ell}\hat{P}_k}$$

$$\hat{U}_j^{\ell}(\theta_j^{\ell}) = \prod_k \left[\cos(\theta_j^{\ell}\beta_k^{j\ell})\hat{I} - i\sin(\theta_j^{\ell}\beta_k^{j\ell})\hat{P}_k \right]$$

Theory: QNN's & Gradients

- Quantum Expectation estimation

$$f(\boldsymbol{\theta}) = \langle \hat{H} \rangle_{\boldsymbol{\theta}} \equiv \langle \Psi_0 | \hat{U}^\dagger(\boldsymbol{\theta}) \hat{H} \hat{U}(\boldsymbol{\theta}) | \Psi_0 \rangle$$

$$\hat{H} = \sum_{k=1}^N \alpha_k \hat{h}_k \equiv \boldsymbol{\alpha} \cdot \hat{\mathbf{h}}, \quad \hat{h}_j = \sum_{k=1}^{J_j} \gamma_k^j \hat{P}_k,$$

$$f(\boldsymbol{\theta}) = \langle \hat{H} \rangle_{\boldsymbol{\theta}} = \sum_{k=1}^N \alpha_k \langle \hat{h}_k \rangle_{\boldsymbol{\theta}} \equiv \boldsymbol{\alpha} \cdot \mathbf{h}_{\boldsymbol{\theta}},$$

- Can estimate expectation values of observables by sampling the output of the quantum computer in various bases

Theory: QNN's & Gradients

- Finite-difference gradients

$$f(\boldsymbol{\theta}) = \langle \hat{H} \rangle_{\boldsymbol{\theta}} \equiv \langle \Psi_0 | \hat{U}^\dagger(\boldsymbol{\theta}) \hat{H} \hat{U}(\boldsymbol{\theta}) | \Psi_0 \rangle$$

$$\partial_k f(\boldsymbol{\theta}) = \frac{f(\boldsymbol{\theta} + \varepsilon \boldsymbol{\Delta}_k) - f(\boldsymbol{\theta} - \varepsilon \boldsymbol{\Delta}_k)}{2\varepsilon} + \mathcal{O}(\varepsilon^2)$$

- Better ways to find gradients?
 - Use chain rule + the fact that QNN's are expressible as composition of exponentials
 - [see TFQ whitepaper](#)

Theory: QNN's & Gradients

see [TFQ whitepaper](#)

- Parameter shift rule $\hat{U}^\ell(\boldsymbol{\theta}^\ell) \equiv \bigotimes_{j=1}^{M_\ell} \hat{U}_j^\ell(\theta_j^\ell)$ $\hat{U}_j^\ell(\theta_j^\ell) = \prod_k e^{-i\theta_j^\ell \beta_k^{j\ell} \hat{P}_k}$

$$\hat{U}^\ell(\boldsymbol{\theta}^\ell) \mapsto \hat{U}^\ell(\boldsymbol{\eta}^\ell) \equiv \bigotimes_{j \in \mathcal{I}_\ell} \left(\prod_k e^{-i\eta_k^{j\ell} \hat{P}_k} \right)$$

$$\eta_k^{j\ell} \equiv \theta_j^\ell \beta_k^{j\ell} \quad f(\boldsymbol{\eta}) \equiv \langle \Psi_0 | \hat{U}^\dagger(\boldsymbol{\eta}) \hat{H} \hat{U}(\boldsymbol{\eta}) | \Psi_0 \rangle$$

$$\frac{\partial f}{\partial \theta_j^\ell} = \sum_k \frac{\partial f}{\partial \eta_k^{j\ell}} \frac{\partial \eta_k^{j\ell}}{\partial \theta_j^\ell} = \sum_k \beta_k^{j\ell} \frac{\partial f}{\partial \eta_k^{j\ell}}$$

$$e^{-i\eta_k^{j\ell} \hat{P}_k} = \cos(\eta_k^{j\ell}) \hat{I} - i \sin(\eta_k^{j\ell}) \hat{P}_k$$

$$\frac{\partial}{\partial \eta_k^{j\ell}} f(\boldsymbol{\eta}) = f(\boldsymbol{\eta} + \frac{\pi}{4} \boldsymbol{\Delta}_k^{j\ell}) - f(\boldsymbol{\eta} - \frac{\pi}{4} \boldsymbol{\Delta}_k^{j\ell})$$

Theory: QNN's & Gradients

see [TFQ whitepaper](#)

- Stochastic gradient estimation

$$\frac{\partial f}{\partial \theta_j^\ell} = \sum_{k=1}^{K_{j\ell}} \beta_k^{j\ell} \frac{\partial f}{\partial \eta_k} = \sum_{k=1}^{K_{j\ell}} \left[\sum_{\pm} \pm \beta_k^{j\ell} f(\boldsymbol{\eta} \pm \frac{\pi}{4} \boldsymbol{\Delta}_k^{j\ell}) \right] \quad k \sim \Pr(k|j, \ell) = |\beta_k^{j\ell}| / (\sum_{o=1}^{K_{j\ell}} |\beta_o^{j\ell}|)$$

- Doubly stochastic gradient estimation

$$f(\boldsymbol{\theta}) = \langle \hat{H} \rangle_{\boldsymbol{\theta}} = \sum_{m=1}^N \alpha_m \langle \hat{h}_m \rangle_{\boldsymbol{\theta}} = \sum_{m=1}^N \sum_{q=1}^{J_m} \alpha_m \gamma_q^m \langle \hat{P}_{qm} \rangle_{\boldsymbol{\theta}},$$
$$\{q, m\} \sim \Pr(q, m) = |\alpha_m \gamma_q^m| / (\sum_{d=1}^N \sum_{r=1}^{J_d} |\alpha_d \gamma_r^d|)$$

- Triply stochastic gradient descent

$$\theta_j^\ell \sim \Pr(j, \ell) = \sum_{k=1}^{K_{j\ell}} |\beta_k^{j\ell}| / (\sum_{u=1}^L \sum_{i=1}^{M_u} \sum_{o=1}^{K_{iu}} |\beta_o^{iu}|)$$

$$\{j, \ell, k, q, m\} \sim \Pr(j, \ell, k, q, m) = \Pr(k|j, \ell) \Pr(j, \ell) \Pr(q, m)$$

Example: a simple quantum state classifier

- App: classifying input quantum states
- Data: [Notebook link](#)
- Model:

Example: a naive MNIST quantum classifier

- App: classifying input classical data encoded in quantum states

- Data:

- Model:

[Notebook link](#)

How to more practically leverage
quantum computing power?

Hybridize Classical & Quantum Deep Learning!

Hybrid Quantum-classical neural networks

Combine quantum and classical representation power

Hybrid Quantum-classical neural networks

- Practical approach to add quantum components to classical deep learning representations
- Can make hybrid computational graphs from new building blocks

Hybrid quantum-classical computational graph

QNN: quantum
neural network

DNN: (Classical)
Deep Neural
Network

Quantum-classical Hybrid neural networks & hybrid backprop

- Feedforward expectation value vector: $(\mathbf{h}_\theta)_k = \langle \hat{h}_k \rangle_\theta \equiv \langle \Psi_0 | \hat{U}^\dagger(\theta) \hat{h}_k \hat{U}(\theta) | \Psi_0 \rangle$

- Effective backpropagated error Hamiltonian:
 - Gradients can be estimated with gradient methods from last lecture

$$\hat{H}_g \equiv \sum_k g_k \hat{h}_k$$

$$\frac{\partial}{\partial \theta_j} \langle \hat{H}_g \rangle_\theta = \frac{\partial}{\partial \theta_j} (\mathbf{g} \cdot \mathbf{h}_\theta) = \sum_k g_k \frac{\partial h_{\theta,k}}{\partial \theta_j}$$

Reminder :Types of Machine Learning

ML: **Algorithms** to identify patterns in data

Generative
models

Discriminative
models

Dataset

$$\mathcal{D} = \{\mathbf{x}_k\}_{k=1}^{|\mathcal{D}|} \quad \mathbf{x}_k \sim p_{\text{true}}(\mathbf{x})$$

Want to learn:

$$q_{\Phi}(\mathbf{x}) \approx p_{\text{true}}(\mathbf{x})$$

Dataset

$$\mathcal{D} = \{\mathbf{x}_k, \mathbf{y}_k\}_k$$

Want to learn:

$$q_{\Phi}(\mathbf{y}|\mathbf{x}) \approx p_{\text{true}}(\mathbf{y}|\mathbf{x})$$

$$q_{\Phi}(\mathbf{y}|\mathbf{x}) = \delta(y - f_{\Phi}(\mathbf{x}))$$

Types of Hybrid Quantum-Classical Deep Learning Models

Generative:

e.g. quantum-probabilistic
generative models

Discriminative:

e.g. quantum-classical
feedforward neural network
classifiers

Hybrid Quantum-Classical (HQC)

Models

+

Software

HQC neural network models:

Breaking up the task of quantum representation learning into quantum & classical components

TensorFlow Quantum:

Integrating quantum capabilities into SOTA ML workflows, with a focus on scalability & performance

Hybrid Quantum-Classical (HQC)

Models

HQC neural network models:

Breaking up the task of quantum representation learning into quantum & classical components

+

Software

TensorFlow Quantum:

Integrating quantum capabilities into SOTA ML workflows, with a focus on scalability & performance

TensorFlow Quantum

an open source library for the rapid prototyping of hybrid quantum-classical models for classical or quantum data.

Some key features:

- Deep hybridization of TF and Cirq
- Automated expectation & gradient calculations
- Parallelizable high-performance simulation with TFQ-qSim
- Accelerating any quantum variational algo workflow

TensorFlow Quantum

an open source library for the rapid prototyping of hybrid quantum-classical models for classical or quantum data.

Some key features:

- Focused on aiding the construction & training of *hybrid quantum-classical models*
- Focused on supporting quantum data (*standard quantum datasets coming soon*)

TensorFlow Quantum

an open source library for the rapid prototyping of hybrid quantum-classical models for classical or quantum data.

Some key features:

- Hybrid quantum-classical backprop
- Choices of differentiators
 - Finite-difference
 - Parameter shift
 - Stochastic parameter shift
 - Adjoint diff (*new*)

$$(\mathbf{h}_{\boldsymbol{\theta}})_k = \langle \hat{h}_k \rangle_{\boldsymbol{\theta}} \equiv \langle \Psi_0 | \hat{U}^\dagger(\boldsymbol{\theta}) \hat{h}_k \hat{U}(\boldsymbol{\theta}) | \Psi_0 \rangle$$

$$\hat{H}_{\mathbf{g}} \equiv \sum_k g_k \hat{h}_k \quad \frac{\partial}{\partial \theta_j} \langle \hat{H}_{\mathbf{g}} \rangle_{\boldsymbol{\theta}} = \frac{\partial}{\partial \theta_j} (\mathbf{g} \cdot \mathbf{h}_{\boldsymbol{\theta}}) = \sum_k g_k \frac{\partial h_{\boldsymbol{\theta},k}}{\partial \theta_j}$$

TensorFlow Quantum

an open source library for the rapid prototyping of hybrid quantum-classical models for classical or quantum data.

Some key features:

- Deep integration into TF
 - allows for high-performance interfacing with other TF-based frameworks ([e.g. TF Probability & TensorBoard](#))

References

Be sure to check out
[tensorflow.org/quantum](https://www.tensorflow.org/quantum)

[\[arXiv:2003.02989\]](https://arxiv.org/abs/2003.02989)


[\[arXiv:1910.02071\]](https://arxiv.org/abs/1910.02071)

Novel hybrid approach:
Quantum-probabilistic Hybrid DL

Pure Quantum Neural Network Model

Previous solutions

Previous quantum neural nets process strictly pure states



Initial state

$$|\xi_0\rangle$$


Output state

$$|\xi_\phi\rangle = \hat{U}(\phi) |\xi_0\rangle$$

Quantum-probabilistic Hybrid Models

Novel solution:

Combining classical probabilistic inference with quantum neural nets



Latent representation

$$\hat{\rho}_\theta = \sum_{x \in \Omega} p_\theta(x) |x\rangle\langle x|$$

Visible representation

$$\hat{\rho}_{\theta\phi} = \hat{U}(\phi) \hat{\rho}_\theta \hat{U}^\dagger(\phi)$$


So what can one do with these?

Applications of Hybrid Quantum-probabilistic models

Variational Quantum Thermalization:

Variationally learning how to create the thermal state given a Hamiltonian


- minimize quantum free energy
- \exists parameter-shift gradients for θ & ϕ



Quantum Modular Hamiltonian learning:

Learning to replicate mixed data as a parameterized thermal state

- minimize quantum cross-entropy
- \exists parameter-shift gradients for θ & ϕ

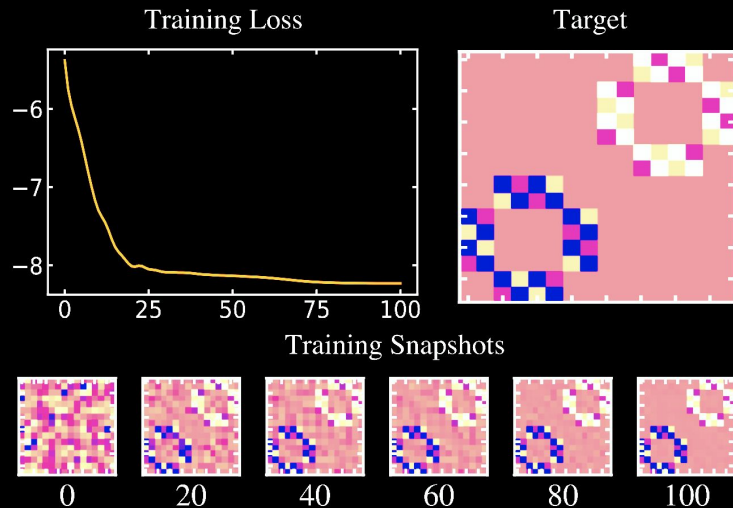


Variational Quantum Thermalization Results

Simulating toy models of superconductors

d-wave superconductor
fermionic Hamiltonian

$$\hat{H}_{d_{x^2-y^2}} = -t \sum_{\langle i,j \rangle, \sigma} (\hat{a}_{i,\sigma}^\dagger \hat{a}_{j,\sigma} + \hat{a}_{j,\sigma}^\dagger \hat{a}_{i,\sigma}) \\ \Delta \sum_{\langle i,j \rangle} (\hat{a}_{i,\uparrow}^\dagger \hat{a}_{j,\downarrow}^\dagger - \hat{a}_{i,\downarrow}^\dagger \hat{a}_{j,\uparrow}^\dagger + \text{h.c.})$$



Learning to generate a thermal state
(covariance matrix plot)

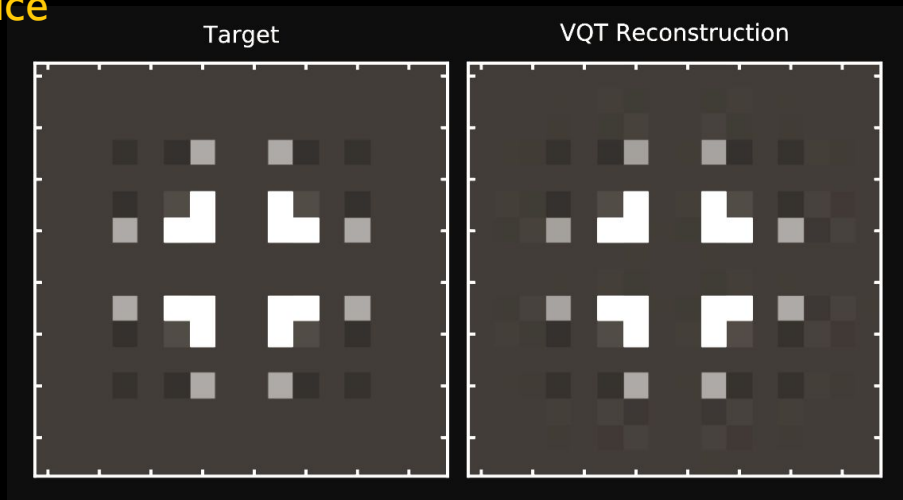
Variational Quantum Thermalization Results

Generating thermal states of 2D spin lattice

2D Heisenberg model

$$\hat{H}_{\text{HEIS}} = \sum_{\langle ij \rangle_h} J_h \hat{\mathbf{S}}_i \cdot \hat{\mathbf{S}}_j + \sum_{\langle ij \rangle_v} J_v \hat{\mathbf{S}}_i \cdot \hat{\mathbf{S}}_j$$

Latent model: product of Bernoullis



Generated thermal state
(density matrix)

[after 200 training steps.]

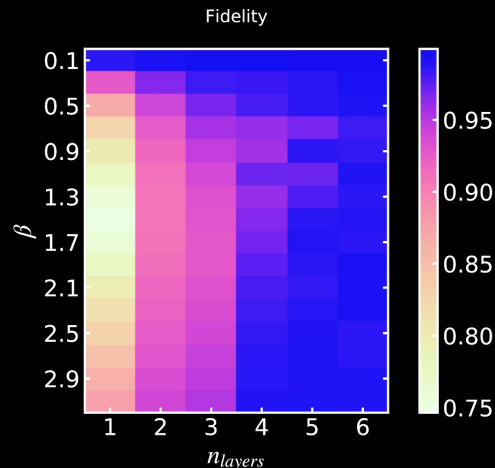
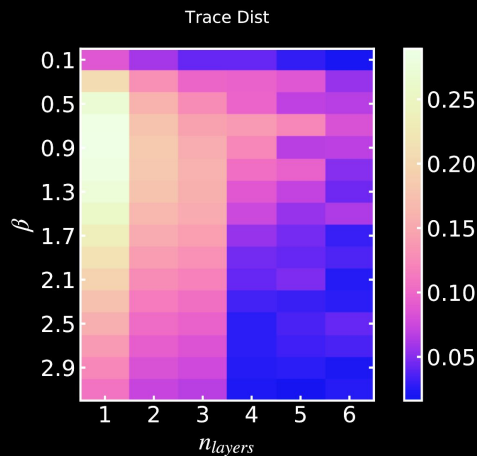

$N_x = N_y = 2$, at $\beta = 2.6$, $J_x = 1.0$, $J_y = 0.6$

Variational Quantum Thermalization

More recent Results

Sweeping over temperatures and QNN circuit depth for the Heisenberg model

Showing a relation between temperature, quantum correlations, and necessary quantum circuit depth



Final trace distance & fidelity after full VQT training

Quantum Variational Thermalization

Preparing Quantum Thermal States with Quantum-probabilistic inference

Quantum Simulation - the Variational Quantum Thermalizer (VQT)

Task: given target \hat{H} and $\beta = 1/T$, generate $\hat{\sigma}_\beta = \frac{1}{\mathcal{Z}_\beta} e^{-\beta \hat{H}}$, $\mathcal{Z}_\beta = \text{tr}(e^{-\beta \hat{H}})$


Preparing Quantum Thermal States with Quantum-probabilistic inference

Quantum Simulation - the Variational Quantum Thermalizer (VQT)

Task: given target \hat{H} and $\beta = 1/T$, generate $\hat{\sigma}_\beta = \frac{1}{\mathcal{Z}_\beta} e^{-\beta \hat{H}}$, $\mathcal{Z}_\beta = \text{tr}(e^{-\beta \hat{H}})$

Variational hybrid model

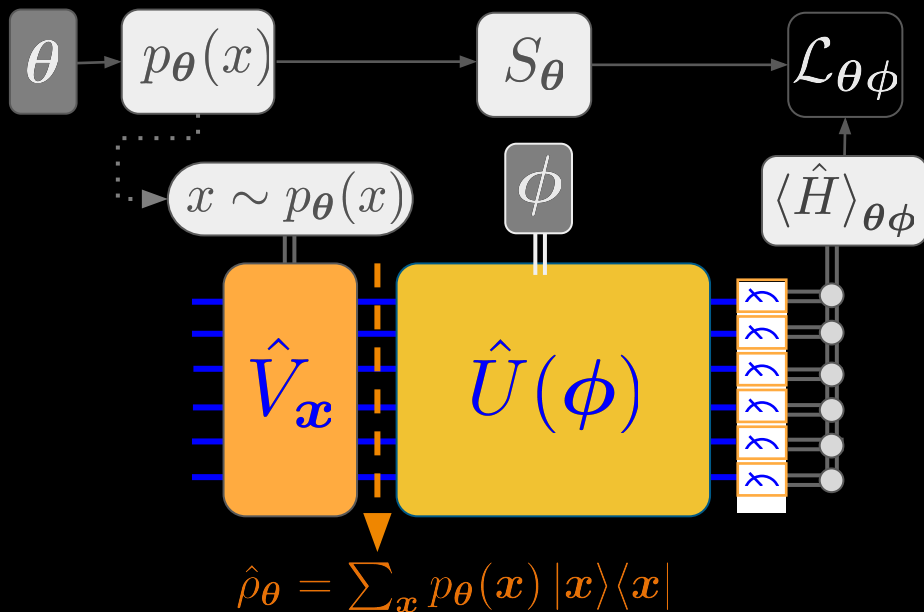
$$\hat{\rho}_{\theta\phi} = \hat{U}(\phi) \hat{\rho}_\theta \hat{U}^\dagger(\phi)$$



Preparing Quantum Thermal States with Quantum-probabilistic inference

Quantum Simulation - the Variational Quantum Thermalizer (VQT)

Task: given target \hat{H} and $\beta = 1/T$, generate $\hat{\sigma}_\beta = \frac{1}{\mathcal{Z}_\beta} e^{-\beta \hat{H}}$, $\mathcal{Z}_\beta = \text{tr}(e^{-\beta \hat{H}})$



Variational hybrid model

$$\hat{\rho}_{\theta\phi} = \hat{U}(\phi) \hat{\rho}_\theta \hat{U}^\dagger(\phi)$$


Minimize free energy

$$\mathcal{L}_{\text{VQT}}(\theta, \phi) = \beta \text{tr}(\hat{\rho}_{\theta\phi} \hat{H}) - S(\hat{\rho}_{\theta\phi})$$

Preparing Quantum Thermal States with Quantum-probabilistic inference

Quantum Simulation - the Variational Quantum Thermalizer (VQT)

Task: given target \hat{H} and $\beta = 1/T$, generate $\hat{\sigma}_\beta = \frac{1}{\mathcal{Z}_\beta} e^{-\beta \hat{H}}$, $\mathcal{Z}_\beta = \text{tr}(e^{-\beta \hat{H}})$



Variational hybrid model

$$\hat{\rho}_{\theta\phi} = \hat{U}(\phi) \hat{\rho}_\theta \hat{U}^\dagger(\phi)$$

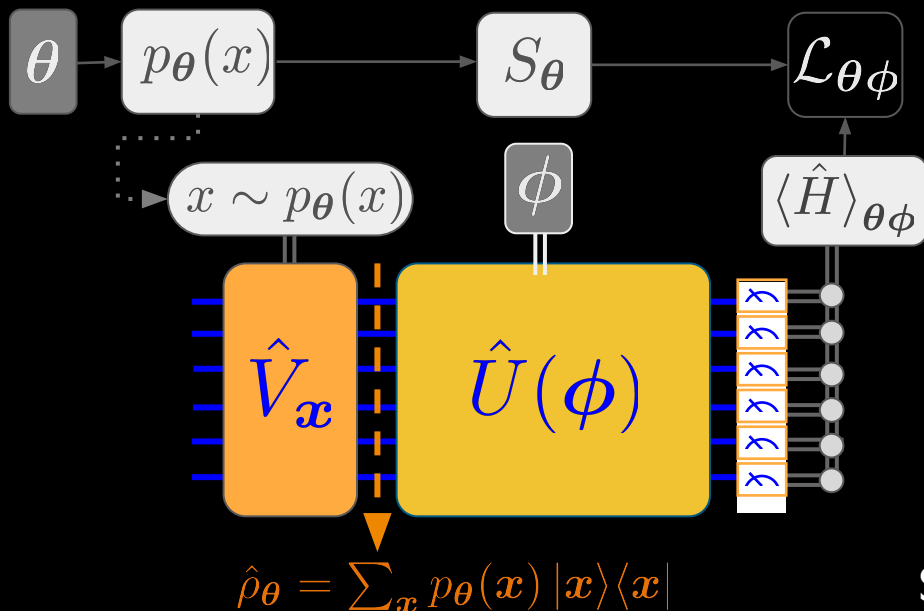
Minimize free energy

$$\mathcal{L}_{\text{VQT}}(\theta, \phi) = \beta \text{tr}(\hat{\rho}_{\theta\phi} \hat{H}) - S(\hat{\rho}_{\theta\phi})$$

Preparing Quantum Thermal States with Quantum-probabilistic inference

Quantum Simulation - the Variational Quantum Thermalizer (VQT)

Task: given target \hat{H} and $\beta = 1/T$, generate $\hat{\sigma}_\beta = \frac{1}{\mathcal{Z}_\beta} e^{-\beta \hat{H}}$, $\mathcal{Z}_\beta = \text{tr}(e^{-\beta \hat{H}})$



Variational hybrid model

$$\hat{\rho}_{\theta\phi} = \hat{U}(\phi) \hat{\rho}_\theta \hat{U}^\dagger(\phi)$$

Minimize free energy

$$\mathcal{L}_{\text{VQT}}(\theta, \phi) = \beta \text{tr}(\hat{\rho}_{\theta\phi} \hat{H}) - S(\hat{\rho}_{\theta\phi})$$

Equivalent to finding

$$\arg\min_{\theta, \phi} D(\hat{\rho}_{\theta\phi} \| \hat{\sigma}_\beta)$$


State of minimal free energy is thermal state!

Quantum Hamiltonian-Based Models

Combining classical probabilistic inference with quantum neural nets

Classical latent distribution
parameterized by classical neural net

$$p_{\theta}(x) = \frac{1}{\mathcal{Z}_{\theta}} e^{-E_{\theta}(x)}$$



Latent Modular Hamiltonian representation

$$\hat{K}_{\theta} = \sum_{x \in \Omega} E_{\theta}(x) |x\rangle\langle x|$$

$$\hat{\rho}_{\theta} = \frac{1}{\mathcal{Z}_{\theta}} e^{-\hat{K}_{\theta}}, \quad \mathcal{Z}_{\theta} = \text{tr}[e^{-\hat{K}_{\theta}}]$$

Visible representation

$$\hat{\rho}_{\theta\phi} = \frac{1}{\mathcal{Z}_{\theta}} e^{-\hat{U}(\phi) \hat{K}_{\theta} \hat{U}^{\dagger}(\phi)} \equiv \frac{1}{\mathcal{Z}_{\theta}} e^{-\hat{K}_{\theta\phi}}$$

Generative Learning of Quantum Mixed States with Quantum Hamiltonian-Based Models


Quantum Modular Hamiltonian Learning for generative modelling

Task: given unknown $\hat{\sigma}_{\mathcal{D}} = \sum_{d \in \mathcal{D}} p_d \hat{\sigma}_d$ find $\{\theta^*, \varphi^*\}$ such that $\hat{\rho}_{\theta^*, \varphi^*} \approx \hat{\sigma}_{\mathcal{D}}$

Generative Learning of Quantum Mixed States with Quantum Hamiltonian-Based Models

Quantum Modular Hamiltonian Learning for generative modelling

Task: given unknown $\hat{\sigma}_{\mathcal{D}} = \sum_{d \in \mathcal{D}} p_d \hat{\sigma}_d$ find $\{\theta^*, \varphi^*\}$ such that $\hat{\rho}_{\theta^* \phi^*} \approx \hat{\sigma}_{\mathcal{D}}$



Approach:

Quantum Hamiltonian-Based Model:

$$\hat{\rho}_{\theta \phi} = \frac{1}{\mathcal{Z}_\theta} e^{-\hat{U}(\phi) \hat{K}_\theta \hat{U}^\dagger(\phi)} \equiv \frac{1}{\mathcal{Z}_\theta} e^{-\hat{K}_{\theta \phi}}$$


Minimize cross-entropy:

$$\mathcal{L}_{\text{QMLH}}(\theta, \phi) \equiv -\text{tr}(\hat{\sigma}_{\mathcal{D}} \log \hat{\rho}_{\theta \phi}) = \text{tr}(\hat{\sigma}_{\mathcal{D}} \hat{K}_{\theta \phi}) + \log(\mathcal{Z}_\theta)$$

Generative Learning of Quantum Mixed States with Quantum Hamiltonian-Based Models

Quantum Modular Hamiltonian Learning for generative modelling

Task: given unknown $\hat{\sigma}_{\mathcal{D}} = \sum_{d \in \mathcal{D}} p_d \hat{\sigma}_d$ find $\{\theta^*, \phi^*\}$ such that $\hat{\rho}_{\theta^* \phi^*} \approx \hat{\sigma}_{\mathcal{D}}$



Approach:

Quantum Hamiltonian-Based Model:

$$\hat{\rho}_{\theta\phi} = \frac{1}{\mathcal{Z}_\theta} e^{-\hat{U}(\phi) \hat{K}_\theta \hat{U}^\dagger(\phi)} \equiv \frac{1}{\mathcal{Z}_\theta} e^{-\hat{K}_{\theta\phi}}$$

Minimize cross-entropy:

$$\mathcal{L}_{\text{QMH}}(\theta, \phi) \equiv -\text{tr}(\hat{\sigma}_{\mathcal{D}} \log \hat{\rho}_{\theta\phi}) = \text{tr}(\hat{\sigma}_{\mathcal{D}} \hat{K}_{\theta\phi}) + \log(\mathcal{Z}_\theta)$$

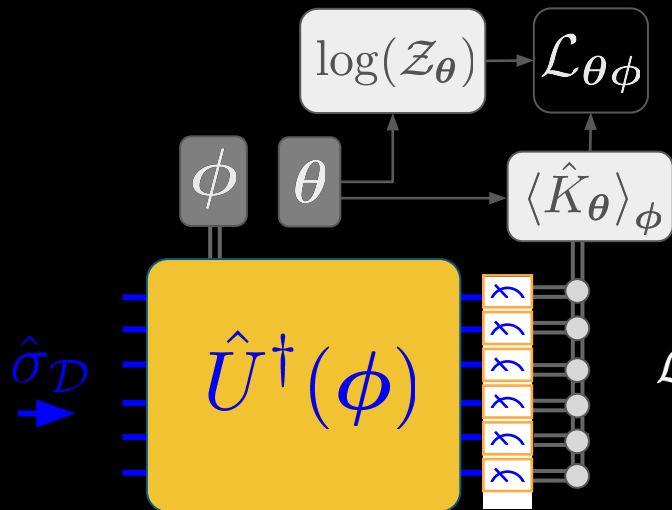
Equiv. to minimizing quantum relative entropy

$$\underset{\theta, \phi}{\operatorname{argmin}} D(\hat{\sigma}_{\mathcal{D}} \| \hat{\rho}_{\theta\phi})$$

Generative Learning of Quantum Mixed States with Quantum Hamiltonian-Based Models

Quantum Modular Hamiltonian Learning for generative modelling

Task: given unknown $\hat{\sigma}_{\mathcal{D}} = \sum_{d \in \mathcal{D}} p_d \hat{\sigma}_d$ find $\{\theta^*, \phi^*\}$ such that $\hat{\rho}_{\theta^* \phi^*} \approx \hat{\sigma}_{\mathcal{D}}$



Approach:

Quantum Hamiltonian-Based Model:

$$\hat{\rho}_{\theta \phi} = \frac{1}{\mathcal{Z}_\theta} e^{-\hat{U}(\phi) \hat{K}_\theta \hat{U}^\dagger(\phi)} \equiv \frac{1}{\mathcal{Z}_\theta} e^{-\hat{K}_{\theta \phi}}$$

Minimize cross-entropy:

$$\mathcal{L}_{\text{QMHL}}(\theta, \phi) \equiv -\text{tr}(\hat{\sigma}_{\mathcal{D}} \log \hat{\rho}_{\theta \phi}) = \text{tr}(\hat{\sigma}_{\mathcal{D}} \hat{K}_{\theta \phi}) + \log(\mathcal{Z}_\theta)$$

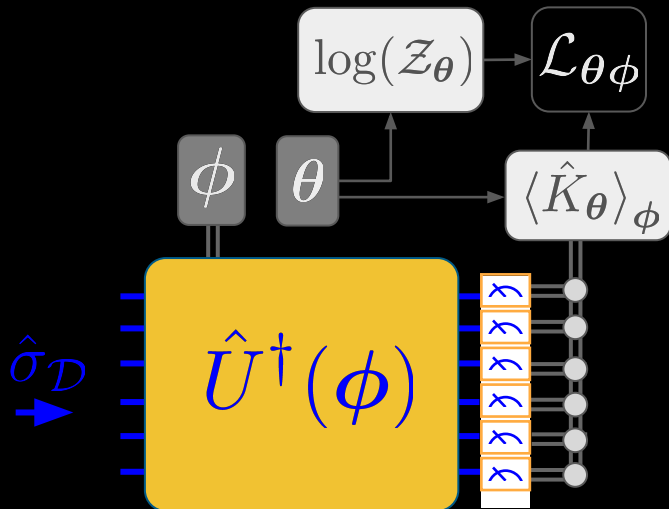
Cross entropy converges to entropy of data

$$\mathcal{L}_{\text{QMHL}}(\theta, \phi) \xrightarrow{\hat{\rho}_{\theta \phi} \rightarrow \hat{\sigma}_{\mathcal{D}}} S(\hat{\sigma}_{\mathcal{D}})$$

Generative Learning of Quantum Mixed States with Quantum Hamiltonian-Based Models

Quantum Modular Hamiltonian Learning for generative modelling

Gradients? $\mathcal{L}_{\text{QMHL}}(\theta, \phi) \equiv -\text{tr}(\hat{\sigma}_{\mathcal{D}} \log \hat{\rho}_{\theta\phi}) = \text{tr}(\hat{\sigma}_{\mathcal{D}} \hat{K}_{\theta\phi}) + \log(\mathcal{Z}_{\theta})$



Classical Model Gradients

$$\begin{aligned}\nabla_{\theta} \mathcal{L}_{\text{QMHL}}(\theta, \phi) &= \nabla_{\theta} \text{tr}([\hat{U}^{\dagger}(\phi) \hat{\sigma}_{\mathcal{D}} \hat{U}(\phi)] \hat{K}_{\theta}) + \nabla_{\theta} \log(\mathcal{Z}_{\theta}) \\ &= \mathbb{E}_{x \sim \sigma_{\phi}(x)} [\nabla_{\theta} E_{\theta}(x)] - \mathbb{E}_{y \sim p_{\theta}(y)} [\nabla_{\theta} E_{\theta}(y)]\end{aligned}$$

$$\text{where } \sigma_{\phi}(x) \equiv \langle x | \hat{U}^{\dagger}(\phi) \hat{\sigma}_{\mathcal{D}} \hat{U}(\phi) | x \rangle$$

QNN gradients: parameter shift on reverse unitary